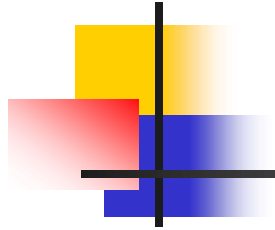




Sorting data with MapReduce

Partial, Total and Secondary Sorting



Partial sorting

Partial sorting

- By default MapReduce results are sorted using the key. The following output refers to the classic “word count”

```
usermr@hadoop: ~/examples/java/Ex01-b
Result sorted by key - MapReduce defaults - (first 5 lines)
hadoop fs -text file:///home/usermr/examples/output/gutenberg/part-r-00000 2>/dev/null | head -n 5
"AS-IS".      8
"Brittain."   2
"Brittish",   1
"Defects".    8
"Defects,"    2

Result sorted (by value) using the linux sort command
hadoop fs -text file:///home/usermr/examples/output/gutenberg/part-r-00000 2>/dev/null | sort -k 2,2 -n -r | head -n 5
the      2700
of       1981
and      1300
to       1282
in       731
usermr@hadoop:~/examples/java/Ex01-b$
```



Partial sorting

- The sort order for keys is controlled by an instance of a `RawComparator`
 1. `mapreduce.job.output.key.comparator.class` property set (explicitly or by calling `setSortComparatorClass()`) is used to create an instance of a `RawComparator`
 2. Else, keys must be a subclass of `WritableComparable`, and the registered comparator for the key class is used
 3. If there is no registered comparator, then a `RawComparator` is used



Partial sorting

- By default, Hadoop assumes a single reducer. This behaviour can be overridden: programmatically `setNumReduceTasks()` or using the property `mapreduce.job.reduces`
- If more than one reducer is used, each output is sorted (by the mappers). However, there is no easy way to combine the files (by concatenation, for example, in the case of plain-text files) to produce a globally sorted file

Partial sorting

usermr@hadoop: ~/examples/java/Ex01-b

Result sorted by key - MapReduce defaults - (first 5 lines)

```
hadoop fs -text file:///home/usermr/examples/output/gutenberg/part-r-00000 2>/dev/null | head -n 5
```

```
"AS-IS".      8
"Brittain."   2
"Brittish",   1
"I            2
"More        1
```

Result sorted (by value) using the linux sort command

```
hadoop fs -text file:///home/usermr/examples/output/gutenberg/part-r-00000 2>/dev/null | sort -k 2,2 -n -r | head -n 5
```

```
the      2700
of       1981
and      1300
to       1282
in       731
```

Result sorted by key - MapReduce defaults - (first 5 lines)

```
hadoop fs -text file:///home/usermr/examples/output/gutenberg/part-r-00001 2>/dev/null | head -n 5
```

```
"Defects".    8
"Defects,"    2
"House"       2
"Information"  2
"Plain"       4
```

Result sorted (by value) using the linux sort command

```
hadoop fs -text file:///home/usermr/examples/output/gutenberg/part-r-00001 2>/dev/null | sort -k 2,2 -n -r | head -n 5
```

```
shall     470
this      402
is        334
any       330
as        258
```

usermr@hadoop:~/examples/java/Ex01-b\$



Partial sorting

- The following example shows a MapReduce application capable of filter the input used in the Temperatures examples
- The example removes from the input dataset all the records that do not contain a valid temperature
- The output is stored in sequence files in the format:
 - `<Temperature (float)><Temperature record (Text)>`
- The number of reducers is zero (no reducers used)

Partial sorting

Ex05

```
public class CleanerMapper
    extends Mapper<LongWritable, Text, FloatWritable, Text> {

    private FloatWritable theKey = new FloatWritable();
    private NcdcRecordParser p = new NcdcRecordParser();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        p.parse( value );
        if ( p.isValidTemperature() && p.getAirTemperature()<100.0F ){
            theKey.set( p.getAirTemperature() );
            context.write( theKey, value );
        }
    }
}
```


Partial sorting

Ex05

```
public class CleanerApplication extends Configured implements Tool {

    public static void main(String[] args)
        throws Exception {
        int exitCode = ToolRunner.run( new CleanerApplication(), args);
        System.exit( exitCode );
    }

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            ...
            return -1;
        }
    }
}
```

Partial sorting

Ex05

```
Job job = Job.getInstance( getConf() );

job.setJobName( "Max temperature - Version 5" );
job.setJarByClass( CleanerApplication.class );

FileInputFormat.addInputPath(job, new Path(args[0]) );
FileOutputFormat.setOutputPath(job, new Path(args[1]) );

job.setMapperClass( CleanerMapper.class );
job.setMapOutputKeyClass( FloatWritable.class );
job.setMapOutputValueClass( Text.class );

job.setNumReduceTasks( 0 );
job.setOutputKeyClass( FloatWritable.class );
job.setOutputValueClass( Text.class );
...
```

Partial sorting

Ex05

```
// -Dmapreduce.job.outputformat.class=...
job.setOutputFormatClass( SequenceFileOutputFormat.class );

// -Dmapreduce.output.fileoutputformat.compress=...
SequenceFileOutputFormat.setCompressOutput( job, true);

// -Dmapreduce.output.fileoutputformat.compress.codec=...
SequenceFileOutputFormat.setOutputCompressorClass( job,
GzipCodec.class );

// -Dmapreduce.output.fileoutputformat.compress.type=...
SequenceFileOutputFormat.setOutputCompressionType( job,
CompressionType.BLOCK );

return job.waitForCompletion(true) ? 0 : 1;
}
}
```

usermr@hadoop: ~/examples/java/Ex01-a5

```
Contents of file:///home/usermr/examples/output/temperatures/small/part-m-00000 (1093 entries):  
-10.6 0029029810999991901010106004+59500+020350FM-12+002699999V0200501N00211999999N0000001N9-01061+99999102641ADDGF10299199999  
-7.2 0029029810999991901010113004+59500+020350FM-12+002699999V0202701N00211999999N0000001N9-00721+99999102671ADDGF10599199999  
-8.3 0029029810999991901010120004+59500+020350FM-12+002699999V0202701N00721999999N0000001N9-00831+99999102671ADDGF10599199999  
...  
2.8 0029029810999991901123106004+59500+020350FM-12+002699999V0201801N00821999999N0000001N9+00281+99999100241ADDGF10899199999  
2.8 0029029810999991901123113004+59500+020350FM-12+002699999V0201601N01181999999N0000001N9+00281+99999100031ADDGF10899199999  
2.2 0029029810999991901123120004+59500+020350FM-12+002699999V0201401N01181999999N0000001N9+00221+99999099681ADDGF10899199999
```

```
Contents of file:///home/usermr/examples/output/temperatures/small/part-m-00001 (1051 entries):  
-5.0 0029029810999991903010106004+59500+020350FM-12+002699999V0200501N01901999999N0000001N9-00501+99999099441ADDGF10899199999  
-5.0 0029029810999991903010113004+59500+020350FM-12+002699999V0200501N01901999999N0000001N9-00501+99999099241ADDGF10899199999  
-2.8 0029029810999991903010120004+59500+020350FM-12+002699999V0200501N01901999999N0000001N9-00281+99999099999  
...  
1.7 0029029810999991903123106004+59500+020350FM-12+002699999V0203601N00511999999N0000001N9+00171+99999101111ADDGF10899199999  
1.1 0029029810999991903123113004+59500+020350FM-12+002699999V0203601N00511999999N0000001N9+00111+99999101111ADDGF10899199999  
1.1 0029029810999991903123120004+59500+020350FM-12+002699999V0203601N00511999999N0000001N9+00111+99999101111ADDGF10899199999
```

Sorted within each file,
but not globally sorted

Equivalent
sizes

```
Contents of file:///home/usermr/examples/output/temperatures/small/part-m-00004 (1098 entries):  
-20.0 0029228920999991908010106004+60717+028783FM-12+000399999V0202701N00101999999N0000001N9-02001+99999103021ADDGF10899199999  
-18.9 0029228920999991908010113004+60717+028783FM-12+000399999V0209991C00001999999N0000001N9-01891+99999103021ADDGF10899199999  
-17.2 0029228920999991908010120004+60717+028783FM-12+000399999V0209991C00001999999N0000001N9-01721+99999102961ADDGF10899199999  
...  
-6.7 0029228920999991908123106004+60717+028783FM-12+000399999V0209991C00001999999N0000001N9-00671+99999104321ADDGF10899199999  
-5.6 0035228920999991908123113004+60717+028783FM-12+000399999V0202501N00671999999N0000001N9-00561+99999104431ADDGF10899199999  
-6.7 0029228920999991908123120004+60717+028783FM-12+000399999V0202001N00671999999N0000001N9-00671+99999104381ADDGF10899199999
```

```
Contents of file:///home/usermr/examples/output/temperatures/small/part-m-00005 (1082 entries):  
-16.1 0029029070999991903010106004+64333+023450FM-12+000599999V0200901N00411999999N0000001N9-01611+99999100911ADDGF10699199999  
-15.6 0029029070999991903010113004+64333+023450FM-12+000599999V0200901N00621999999N0000001N9-01561+99999100831ADDGF10699199999  
-12.2 0029029070999991903010120004+64333+023450FM-12+000599999V0200501N01181999999N0000001N9-01221+99999100401ADDGF10499199999  
...  
-0.6 0029029070999991903123106004+64333+023450FM-12+000599999V0202701N00621999999N0000001N9-00061+99999101881ADDGF10599199999  
0.6 0029029070999991903123113004+64333+023450FM-12+000599999V0202701N00621999999N0000001N9+00061+99999101601ADDGF10599199999  
1.1 0029029070999991903123120004+64333+023450FM-12+000599999V0202701N01591999999N0000001N9+00111+99999101401ADDGF10099199999
```

Total number of entries: 6515



Partial sorting

- In the previous example we had:
 - No reducers
 - The size of each output is similar (input size was similar)
 - The output of each mapper is sorted
- What happens if we use reducers to process the output of the mappers?

Partial sorting

Ex06

```
public class SortByTemperatureMapper
    extends Mapper<FloatWritable, Text, FloatWritable, Text> {
    @Override
    public void map(FloatWritable key, Text value, Context context)
        throws IOException, InterruptedException {
        context.write( key, value );
    }
}
```

```
-----

public class SortByTemperatureReducer
    extends Reducer<FloatWritable, Text, FloatWritable, Text> {
    @Override
    public void reduce(FloatWritable k, Iterable<Text> vs, Context c)
        throws IOException, InterruptedException {
        for (Text value : vs) {
            context.write(k, value );
        }
    }
}
```



Partial sorting

```
public class SortByTemperatureUsingHashPartitionerApplication
    extends Configured implements Tool {

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(
            new SortByTemperatureUsingHashPartitionerApplication(),
            args);
        System.exit( exitCode );
    }

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            ...
            return -1;
        }
        ...
    }
}
```




Partial sorting

```
...  
Job job = Job.getInstance( getConf() );  
  
job.setJobName( "Max temperature - Version 6" );  
job.setJarByClass( SortByTemperatureUsingHashPartitionerApplication.class );  
  
FileInputFormat.addInputPath(job, new Path(args[0]) );  
FileOutputFormat.setOutputPath(job, new Path(args[1]) );  
  
job.setInputFormatClass( SequenceFileInputFormat.class );  
  
job.setMapperClass( SortByTemperatureMapper.class );  
job.setMapOutputKeyClass( FloatWritable.class );  
job.setMapOutputValueClass( Text.class );  
  
...
```




Partial sorting

```
...
job.setOutputFormatClass( SequenceFileOutputFormat.class );
SequenceFileOutputFormat.setCompressOutput( job, true);
SequenceFileOutputFormat.setOutputCompressorClass( job, GzipCodec.class );

SequenceFileOutputFormat.setOutputCompressionType(
    job, CompressionType.BLOCK );

job.setReducerClass( SortByTemperatureReducer.class );
job.setOutputKeyClass( FloatWritable.class );
job.setOutputValueClass( Text.class );

return job.waitForCompletion(true) ? 0 : 1;
}
}
```

Partial sorting

```
usermr@hadoop: ~/examples/java/Ex01-a6

Contents of file:///home/usermr/examples/output/temperatures/small-seq/part-r-00000 (772 entries):
-25.0 0029029070999991901021506004+64333+023450FM-12+000599999V0201801N00821999999N0000001N9-02501+99999102791ADDGF1009919999999999999999
-25.0 0029029440999991908011013004+61500+023767FM-12+008499999V0200701N00311999999N0000001N9-02501+99999100241ADDGF1009919999999999999999
-20.0 0029228920999991908010920004+60717+028783FM-12+000399999V0200701N00261999999N0000001N9-02001+99999101561ADDGF1009919999999999999999
...
25.0 0029029810999991901072306004+59500+020350FM-12+002699999V0209991C00001999999N0000001N9+02501+99999101441ADDGF1029919999999999999999
25.0 0029029810999991901071813004+59500+020350FM-12+002699999V0202301N00211999999N0000001N9+02501+99999102431ADDGF1069919999999999999999
25.0 0029029810999991901072313004+59500+020350FM-12+002699999V0201401N00211999999N0000001N9+02501+99999101351ADDGF1019919999999999999999

Contents of file:///home/usermr/examples/output/temperatures/small-seq/part-r-00001 (1502 entries):
-21.1 0029029070999991901021420004+64333+023450FM-12+000599999V0209991C00001999999N0000001N9-02111+99999103191ADDGF1009919999999999999999
-21.1 0029029440999991908031306804+61500+023767FM-12+008499999V0200901N00211999999N0000001N9-02721+99999100101ADDGF1009919999999999999999
-21.1 0029029440999991908031406004+61500+023767FM-12+008499999V0200901N00411999999N0000001N9-02671+99999101901ADDGF1089919999999999999999
...
26.1 0029228920999991908073113004+60717+028783FM-12+000399999V0209991C00001999999N0000001N9+02831+99999100351ADDGF1059919999999999999999
26.1 0029228920999991908062013004+60717+028783FM-12+008499999V0202701N00721999999N0000001N9+02831+99999101241ADDGF1059919999999999999999
26.1 0029228920999991908072113004+60717+028783FM-12+000399999V0201601N00931999999N0000001N9+02891+99999100601ADDGF1009919999999999999999

Contents of file:///home/usermr/examples/output/temperatures/small-seq/part-r-00002 (2878 entries):
-27.2 0029029440999991908011020004+61500+023767FM-12+008499999V0200901N00211999999N0000001N9-02721+99999100101ADDGF1009919999999999999999
-26.7 0029029440999991908010106004+61500+023767FM-12+008499999V0200901N00411999999N0000001N9-02671+99999101901ADDGF1089919999999999999999
-23.3 0029029440999991908010820004+61500+023767FM-12+008499999V0209991C00001999999N0000001N9-02331+99999100381ADDGF1089919999999999999999
...
28.3 0035029440999991908071913004+61500+023767FM-12+008499999V0201401N00821999999N0000001N9+02831+99999100351ADDGF1059919999999999999999
28.3 0029029440999991908072913004+61500+023767FM-12+008499999V0202701N00721999999N0000001N9+02831+99999101241ADDGF1059919999999999999999
28.3 0029029440999991908072013004+61500+023767FM-12+008499999V0202701N00721999999N0000001N9+02831+99999101241ADDGF1069919999999999999999

Contents of file:///home/usermr/examples/output/temperatures/small-seq/part-r-00003 (1363 entries):
-28.9 0029228920999991908011106004+60717+028783FM-12+000399999V0209991C00001999999N0000001N9-02831+99999101141ADDGF1009919999999999999999
-24.4 0029029440999991908011006004+61500+023767FM-12+008499999V0200501N00871999999N0000001N9-02441+99999100361ADDGF1049919999999999999999
-23.9 0029029440999991908010113004+61500+023767FM-12+008499999V0202301N00211999999N0000001N9-02391+99999102021ADDGF1089919999999999999999
...
28.9 0029228920999991908072113004+60717+028783FM-12+000399999V0201601N00931999999N0000001N9+02891+99999100601ADDGF1009919999999999999999
28.9 0029029440999991908072813004+61500+023767FM-12+008499999V0202301N00511999999N0000001N9+02891+99999101481ADDGF1039919999999999999999
29.4 0029029810999991901072013004+59500+020350FM-12+002699999V0209991C00001999999N0000001N9+02941+99999102501ADDGF1039919999999999999999

Total number of entries: 6515

usermr@hadoop:~/examples/java/Ex01-a6$
```


Sorted within each file,
but not globally sorted

Different
sizes !!!



Partial sorting

- In the previous example we had:
 - Reducers
 - The size of each input (of mappers) is similar
 - The size of each output (of reducers) is very different
 - The output of each mapper is sorted
 - The output of each reducer is sorted
 - Difficult to produce a global sorted output



Total Sorting



Total sorting

- How can we produce a globally sorted file using Hadoop?
 - Use a single partition:
 - This is incredibly inefficient for large files, because one machine must process all the output
 - Produce a set of sorted files that, if concatenated, form a globally sorted file
 - Use a partitioner that respects the total order of the output
 - Choose the partition sizes carefully to ensure that they are fairly even, so job times are not dominated by a single reducer



Total sorting

- Hadoop allows the sampling of the key space:
 - Look at a small subset of the keys to approximate the key distribution, which is then used to construct partitions
- Hadoop comes with a selection of samplers. The `InputSampler` class defines a nested `Sampler` interface whose implementations return a sample of keys given an `InputFormat` and `Job`:

```
public interface Sampler<K,V> {  
    K[] getSample(InputFormat<K,V> inf, Job job)  
        throws IOException, InterruptedException;  
}
```



Total sorting

- Usually, this interface is not called directly. Instead, the `writePartitionFile()` static method (of class `InputSampler`) is used to create a sequence file that stores the keys that define the partitions

```
public static <K, V> void writePartitionFile(Job j, Sampler<K, V> sampler)
    throws IOException, ClassNotFoundException, InterruptedException {
    ...
}
```

- The sequence file is used by `TotalOrderPartitioner` to create partitions for the sort job.



Total sorting

```
public class SortByTemperatureUsingTotalOrderPartitionerApplication
    extends Configured implements Tool {

    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(
            new SortByTemperatureUsingTotalOrderPartitionerApplication(),
            args);
        System.exit( exitCode );
    }

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            ...
            return -1;
        }
        ...
    }
}
```


Total sorting

```
...
Configuration conf = getConf();
Job job = Job.getInstance( conf );
job.setJobName( "Max temperature - Version 7" );
job.setJarByClass(
SortByTemperatureUsingTotalOrderPartitionerApplication.class );

FileInputFormat.addInputPath(job, new Path(args[0]) );
FileOutputFormat.setOutputPath(job, new Path(args[1]) );

job.setInputFormatClass( SequenceFileInputFormat.class );
job.setMapperClass( SortByTemperatureMapper.class );

job.setMapOutputKeyClass( FloatWritable.class );
job.setMapOutputValueClass( Text.class );
...
```

Total sorting

```
...  
job.setReducerClass( SortByTemperatureReducer.class );  
job.setOutputKeyClass( FloatWritable.class );  
job.setOutputValueClass( Text.class );  
  
job.setOutputFormatClass( SequenceFileOutputFormat.class );  
SequenceFileOutputFormat.setCompressOutput( job, true);  
SequenceFileOutputFormat.setOutputCompressorClass(  
    job, GzipCodec.class );  
SequenceFileOutputFormat.setOutputCompressionType(  
    job, CompressionType.BLOCK );  
  
job.setPartitionerClass( TotalOrderPartitioner.class );  
InputSampler.Sampler<FloatWritable, Text> sampler;  
sampler = new  
    InputSampler.RandomSampler<FloatWritable, Text>(0.1, 10000, 10);  
...
```

Total sorting

```
InputSampler.writePartitionFile( job, sampler );
```

```
String partitionFile;  
partitionFile = TotalOrderPartitioner.getPartitionFile( conf );  
URI partitionUri = new URI( partitionFile );  
job.addCacheFile( partitionUri );
```

```
System.out.printf(  
    "Partition file absolute location: hdfs:///user/%s/%s\n",  
    System.getenv("USER"), partitionUri.getPath() );
```

```
return job.waitForCompletion(true) ? 0 : 1;
```

```
}
```

```
}
```

Total sorting

```
2019-11-28 19:53:00,660 INFO input.FileInputFormat: Total input files to process : 6
2019-11-28 19:53:00,697 INFO zlib.ZlibFactory: Successfully loaded & initialized native-zlib library
2019-11-28 19:53:00,698 INFO compress.CodecPool: Got brand-new decompressor [.gz]
2019-11-28 19:53:00,702 INFO compress.CodecPool: Got brand-new decompressor [.gz]
2019-11-28 19:53:00,702 INFO compress.CodecPool: Got brand-new decompressor [.gz]
2019-11-28 19:53:00,702 INFO compress.CodecPool: Got brand-new decompressor [.gz]
2019-11-28 19:53:00,745 INFO partition.InputSampler: Using 677 samples
2019-11-28 19:53:01,280 INFO compress.CodecPool: Got brand-new compressor [.deflate]
Partition file absolute location: hdfs:///user/usermr/ partition.lst
2019-11-28 19:53:01,492 INFO client.RMProxy: Connecting to ResourceManager at hadoop.lrcd.e.ipl.pt/10.
2019-11-28 19:53:01,766 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/ha
0079
2019-11-28 19:53:01,811 INFO input.FileInputFormat: Total input files to process : 6
2019-11-28 19:53:01,847 INFO mapreduce.JobSubmitter: number of splits:6
2019-11-28 19:53:01,977 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1574173509068_0079
2019-11-28 19:53:01,979 INFO mapreduce.JobSubmitter: Executing with tokens: []
2019-11-28 19:53:02,183 INFO conf.Configuration: resource-types.xml not found
2019-11-28 19:53:02,183 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2019-11-28 19:53:02,239 INFO impl.YarnClientImpl: Submitted application application_1574173509068_0079
2019-11-28 19:53:02,271 INFO mapreduce.Job: The url to track the job: http://hadoop.lrcd.e.ipl.pt:8088
2019-11-28 19:53:02,271 INFO mapreduce.Job: Running job: job_1574173509068_0079
2019-11-28 19:53:08,368 INFO mapreduce.Job: Job job_1574173509068_0079 running in uber mode : false
2019-11-28 19:53:08,369 INFO mapreduce.Job:  map 0% reduce 0%
2019-11-28 19:53:15,478 INFO mapreduce.Job:  map 17% reduce 0%
2019-11-28 19:53:19,507 INFO mapreduce.Job:  map 33% reduce 0%
2019-11-28 19:53:22,528 INFO mapreduce.Job:  map 50% reduce 0%
2019-11-28 19:53:24,552 INFO mapreduce.Job:  map 67% reduce 0%
2019-11-28 19:53:25,562 INFO mapreduce.Job:  map 83% reduce 0%
```



```
usermr@hadoop: ~/examples/java/Ex01-a7

Partition contents:
-1.1 (null)
3.9 (null)
11.1 (null)

Contents of file:///home/usermr/examples/output/temperatures/small-sorted/part-r-00000 (1515 entries):
-28.9 0029228920999991908011106004+60717+028783FM-12+000399999V0209991C000019999999N0000001N9-02891+99999101141ADDGF1009919999999999999999
-27.2 0029029440999991908011020004+61500+023767FM-12+008499999V0200901N002119999999N0000001N9-02721+99999100101ADDGF1009919999999999999999
-26.7 0029029440999991908010106004+61500+023767FM-12+008499999V0200901N004119999999N0000001N9-02671+99999101901ADDGF1089919999999999999999
...
-1.7 0029029810999991903040313004+59500+020350FM-12+002699999V0200201N013919999999N0000001N9-00171+99999100661ADDGF1049919999999999999999
-1.7 0029029810999991903021213004+59500+020350FM-12+002699999V0203201N020119999999N0000001N9-00171+99999109781ADDGF1089919999999999999999
-1.7 0029029810999991903122806004+59500+020350FM-12+002699999V0201801N002119999999N0000001N9-00171+99999102481ADDGF1039919999999999999999

Contents of file:///home/usermr/examples/output/temperatures/small-sorted/part-r-00001 (1839 entries):
-1.1 0029029070999991901031420004+64333+023450FM-12+000599999V0201801N006219999999N0000001N9-00111+99999102941ADDGF1089919999999999999999
-1.1 0029029070999991901111213004+64333+023450FM-12+000599999V0200201N015919999999N0000001N9-00111+9999910201ADDGF1089919999999999999999
-1.1 0035029070999991901030920004+64333+023450FM-12+000599999V0200201N015919999999N0000001N9-00111+9999910201ADDGF1089919999999999999999
...
3.3 0029029810999991903050806004+59500+020350FM-12+002699999V0202301N011819999999N0000001N9+01061+99999100671ADDGF1069919999999999999999
3.3 0029029810999991903050713004+59500+020350FM-12+002699999V0201401N003119999999N0000001N9+01061+99999100481ADDGF1069919999999999999999
3.3 0029029810999991903050706004+59500+020350FM-12+002699999V0202701N012919999999N0000001N9+00391+99999100031ADDGF1079919999999999999999

Contents of file:///home/usermr/examples/output/temperatures/small-sorted/part-r-00002 (1428 entries):
3.9 0029029810999991901012313004+59500+020350FM-12+002699999V0202901N018019999999N0000001N9+00391+999991099741ADDGF1009919999999999999999
3.9 0029029810999991901043020004+59500+020350FM-12+002699999V0200701N007219999999N0000001N9+00391+99999102421ADDGF1019919999999999999999
3.9 0029029810999991901111013004+59500+020350FM-12+002699999V0202701N012919999999N0000001N9+00391+999991099841ADDGF1029919999999999999999
...
10.6 0029029810999991901100920004+59500+020350FM-12+002699999V0202301N011819999999N0000001N9+01061+999991099681ADDGF1069919999999999999999
10.6 0029029810999991901101006004+59500+020350FM-12+002699999V0201401N003119999999N0000001N9+01061+99999100481ADDGF1029919999999999999999
10.6 0029029810999991901101106004+59500+020350FM-12+002699999V0201801N008719999999N0000001N9+01061+99999101701ADDGF1019919999999999999999

Contents of file:///home/usermr/examples/output/temperatures/small-sorted/part-r-00003 (1733 entries):
11.1 0029029070999991903090606004+64333+023450FM-12+000599999V0202001N013919999999N0000001N9+01111+99999102031ADDGF1089919999999999999999
11.1 0029029070999991903071920004+64333+023450FM-12+000599999V0203601N015919999999N0000001N9+01111+999991099871ADDGF1089919999999999999999
11.1 0029029070999991903090720004+64333+023450FM-12+000599999V0202501N002119999999N0000001N9+01111+99999101021ADDGF1089919999999999999999
...
28.9 0029029440999991908072813004+61500+023767FM-12+008499999V0202301N005119999999N0000001N9+02891+99999101481ADDGF1039919999999999999999
28.9 0029228920999991908072113004+60717+028783FM-12+000399999V0201601N009319999999N0000001N9+02891+99999100601ADDGF1009919999999999999999
29.4 0029029810999991901072013004+59500+020350FM-12+002699999V0209991C000019999999N0000001N9+02941+99999102501ADDGF1039919999999999999999

Total number of entries: 6515


usermr@hadoop:~/examples/java/Ex01-a7$
```

Partitions boundaries

Sorted within each file,
and globally sorted

Equivalent
sizes

Total sorting

 usermr@hadoop: ~/exam

Partition contents:

-1.1 (null)
3.9 (null)
11.1 (null)

< -1.1 °C	[-1.1°C, 3.9 °C [[3.9°C, 11.1°C [> 11.1°C
1515/6515 = 23%	1839/6515 = 28%	1428/6515 = 22%	1733/6515 = 27%

- The file that holds the partition boundaries must be shared by all the tasks running in the cluster. This can be achieved using the distributed cache mechanism




Total sorting

- The input dataset determines the best sampler to use:
 - `SplitSampler`, which samples only the first n records in a split, is not so good for sorted data
 - `IntervalSampler` chooses keys at regular intervals through the split and makes a better choice for sorted data
 - `RandomSampler` is a good general-purpose



Total sorting

- `InputSampler` and `TotalOrderPartitioner` allow us to choose the number of partitions, i.e., the number of reducers.
- However, `TotalOrderPartitioner` will work only if the partition boundaries are distinct.
- Choosing a high number of partitions with a small number of keys may lead to collisions



Secondary Sorting



Secondary sorting

- Scenario 1:
 - Sort the temperature readings by year (ascending order) and for each year sort the temperatures by their value (descending order)
 - For each year we choose the first temperature. This is equivalent to the maximum temperature for a particular year



Secondary sorting

- Scenario 2:
 - In a set of videos identify all the peoples that appear in the videos and produce a report that for each people presents:
 - The videos where each people appear
 - For each video the time where the people appear



Secondary sorting

- Scenario 3:

- A set of sensors used to collect scientific data, where we have m sensors that are sampled at regular time intervals t
- For each sensor we want to obtain the values of the sensors ordered by the timestamp

(t_1, m_1, r_{80521})

(t_1, m_2, r_{14209})

(t_1, m_3, r_{76042})

...

(t_2, m_1, r_{21823})

(t_2, m_2, r_{66508})

(t_2, m_3, r_{98347})

$m_1 \rightarrow (t_1, r_{80521})$



Secondary sorting

- MapReduce framework sorts the records by the key before they reach the reducers. For any particular key, however, the values are not sorted.
- In general MapReduce programs are written in such a way that they not to depend on the order in which the values appear to the reduce function.
- However, it is possible to impose an order on the values by sorting and grouping the keys in a particular way




Secondary sorting

- In the previous scenarios, the mechanism used to sort the data is also known (in the context of MapReduce applications) as **Secondary Sorting**
- In order to use secondary sorting we need to create a combined key that depends on the original key but also depends of the values

Secondary sorting

- In the temperature scenario we want to sort the temperatures by the year, and for each year we want to sort the temperature readings





Secondary sorting

- To implement secondary sorting we need to:
 - Make a new **composite key** formed by the natural key and the natural value(s)
 - Implement a sort comparator capable of sorting the **composite key**
 - Implement a partitioner and grouping comparator that can only consider the **key** (not the composite key) for partitioning and sorting purposes

Secondary sorting

Ex08

```
public class TemperaturePair
    implements WritableComparable<TemperaturePair> {
    private IntWritable year;
    private FloatWritable temperature;

    // Getters and setters for the year and the temperature ...

    public TemperaturePair() {
        this.set(new IntWritable(), new FloatWritable() );
    }

    public TemperaturePair(IntWritable year, FloatWritable temperature) {
        this.set( year, temperature);
    }

    public void set(IntWritable year, FloatWritable temperature) {
        this.year = year;
        this.temperature = temperature;
    }

    ...
}
```

Secondary sorting

Ex08

...

@Override

```
public void write(DataOutput out) throws IOException {  
    this.year.write( out );  
    this.temperature.write( out );  
}
```

@Override

```
public void readFields(DataInput in) throws IOException {  
    this.year.readFields( in );  
    this.temperature.readFields( in );  
}
```

@Override

```
public int hashCode() {  
    return this.year.hashCode() * 163 + this.temperature.hashCode();  
}  
...
```

Secondary sorting

Ex08

```
...
@Override
public boolean equals(Object o) {
    if (o instanceof TemperaturePair) {
        TemperaturePair tp = (TemperaturePair)o;
        return year.equals(tp.year)&&temperature.equals(tp.temperature);
    }
    return false;
}

@Override
public String toString() {return this.year + "\t" + this.temperature;}

@Override
public int compareTo(TemperaturePair tp) {
    int cmp = this.year.compareTo( tp.year );
    if (cmp != 0) { return cmp; }
    return this.temperature.compareTo( tp.temperature );
}
}
```

Secondary sorting

Ex08

```
public class MaxTemperatureKeyComparator
    extends WritableComparator {

    protected MaxTemperatureKeyComparator() {
        super(TemperaturePair.class, true);
    }

    @Override
    public int compare(WritableComparable w1, WritableComparable w2) {
        TemperaturePair tp1 = (TemperaturePair)w1;
        TemperaturePair tp2 = (TemperaturePair)w2;

        int cmp = tp1.getYear().compareTo( tp2.getYear() );
        if ( cmp!=0 ) { return cmp; }

        // Reverse order
        return -tp1.getTemperature().compareTo( tp2.getTemperature() );
    }
}
```

Secondary sorting

Ex08

```
public class MaxTemperatureGroupComparator
    extends WritableComparator {

    protected MaxTemperatureGroupComparator() {
        super(TemperaturePair.class, true);
    }

    @Override
    public int compare(WritableComparable w1, WritableComparable w2) {
        TemperaturePair tp1 = (TemperaturePair) w1;
        TemperaturePair tp2 = (TemperaturePair) w2;

        return tp1.getYear().compareTo( tp2.getYear() );
    }
}
```

Secondary sorting

Ex08

```
public class MaxTemperaturePartitioner
    extends Partitioner<TemperaturePair, FloatWritable> {

    @Override
    public int getPartition(
        TemperaturePair key,
        FloatWritable value,
        int numPartitions) {

        return Math.abs( key.getYear().get() * 127) % numPartitions;
    }
}
```

Secondary sorting

Ex08

```
public class MaxTemperatureMapper
    extends Mapper<LongWritable,Text,TemperaturePair,FloatWritable> {

    private TemperaturePair theKey = new TemperaturePair();
    private NcdcRecordParser p = new NcdcRecordParser();

    @Override
    public void map(LongWritable key, Text value, Context context)
        throws IOException, InterruptedException {

        this.p.parse( value );

        if ( p.isValidTemperature() && p.getAirTemperature()<100.0F ) {
            this.theKey.setYear( this.parser.getYear() );
            this.theKey.setTemperature( this.parser.getAirTemperature() );

            context.write( this.theKey, this.theKey.getTemperature() );
        }
    }
}
```

Secondary sorting

Ex08

```
public class MaxTemperatureReducer
    extends Reducer<TemperaturePair,FloatWritable,TemperaturePair,Text>{

    private List<String> valuesAsString = new ArrayList<String>();

    @Override
    protected void reduce(
        TemperaturePair key,
        Iterable<FloatWritable> values,
        Context context)
        throws IOException, InterruptedException {

        this.valuesAsString.clear();

        for (FloatWritable currentTemperature : values) {
            this.valuesAsString.add(0, "" + currentTemperature.get() );
        }

        ...
    }
}
```


Secondary sorting

Ex08

```
...
String _out;
Text out;

_out = "[" + ... + "]";
out = new Text(
    String.format(
        "{%d temperatures, %s}", this.valuesAsString.size(),
        _out) );

context.write( key, out );
}
}
```

Secondary sorting

Ex08

```
public class MaxTemperatureApplication
    extends Configured implements Tool {

    public static void main(String[] args) throws Exception {
        int exCode = ToolRunner.run( new MaxTemperatureApplication(), args);
        System.exit( exCode );
    }

    @Override
    public int run(String[] args) throws Exception {
        if (args.length != 2) {
            ...
            return -1;
        }

        Job job = Job.getInstance( getConf() );

        job.setJobName( "Max temperature - Version 8" );
        job.setJarByClass( MaxTemperatureApplication.class );
        ...
    }
}
```

Secondary sorting

Ex08

```
...
FileInputFormat.addInputPath(job, new Path(args[0]) );
FileOutputFormat.setOutputPath(job, new Path(args[1]) );
job.setMapperClass( MaxTemperatureMapper.class );
job.setMapOutputKeyClass( TemperaturePair.class );
job.setMapOutputValueClass( FloatWritable.class );
job.setPartitionerClass( MaxTemperaturePartitioner.class );
job.setSortComparatorClass( MaxTemperatureKeyComparator.class );
job.setGroupingComparatorClass( MaxTemperatureGroupComparator.class );
job.setReducerClass( MaxTemperatureReducer.class );
job.setOutputKeyClass( TemperaturePair.class );
job.setOutputValueClass( Text.class );
job.setOutputFormatClass( SequenceFileOutputFormat.class );
SequenceFileOutputFormat.setCompressOutput( job, true );
SequenceFileOutputFormat.setOutputCompressorClass(
    job, GzipCodec.class );
SequenceFileOutputFormat.setOutputCompressionType(
    job, CompressionType.BLOCK );
return job.waitForCompletion(true) ? 0 : 1;
}
}
```

Secondary sorting

```
usermr@hadoop: ~/examples/java/Ex01-a8

Contents of file:///home/usermr/examples/output/temperatures/small/part-r-00000 (1 entries):
1908    -28.9    {2194 temperatures, [-28.9, -27.2, ... , 2.8, ... , 28.9, 28.9]}

Contents of file:///home/usermr/examples/output/temperatures/small/part-r-00001 (2 entries):
1901    -25.0    {2188 temperatures, [-25.0, -21.1, ... , 3.3, ... , 27.2, 29.4]}
1903    -19.4    {2133 temperatures, [-19.4, -18.3, ... , 3.3, ... , 21.7, 21.7]}

Total number of entries: 3

usermr@hadoop:~/examples/java/Ex01-a8$
```

■ $2194 + 2188 + 2133 = 6515$